

Kubernetes

Thorsten Scherf
Red Hat

SLAC 2015, Berlin



Agenda

- What is Kubernetes
- Kubernetes Architecture
- Configuration and Cluster Setup
- Demo

DOCKER PROS AND CONS

- Pros

- Ease of use, makes Linux containers consumable
- Very easy to create and work with derivative images
- Fast boot on containers

- Cons

- Host centric solution, not aware of anything else
- Cannot handle networking between containers on separate hosts
- No higher level provisioning to connect related containers



Kubernetes is not limited to Docker, support for AppC based containers (like rkt) and other application containers is coming (*).

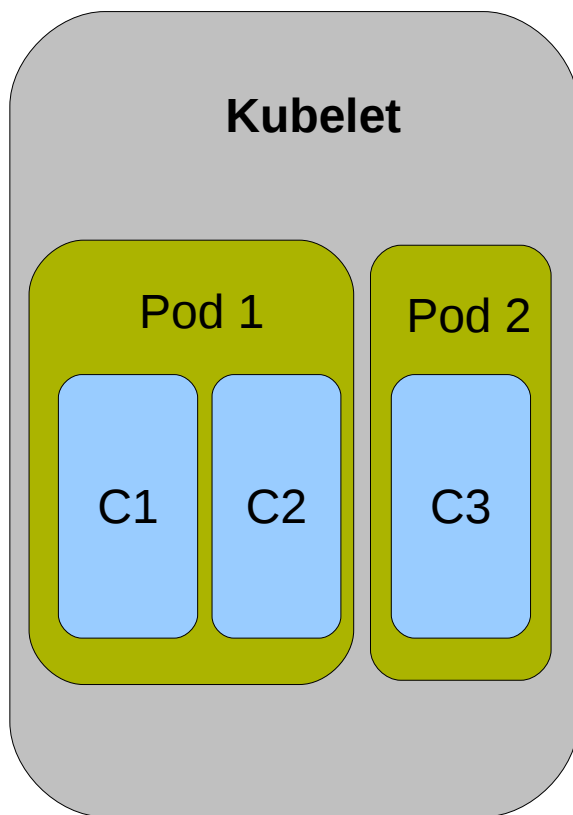
(*) <https://xkcd.com/927/>

What is Kubernetes?

- A highly collaborative open source project originally conceived by Google
 - Google has 10+ years experience w/ containerized apps
- Start, stop, update, and manage a cluster of machines running containers in a consistent and maintainable way.
- Sometimes called:
 - kube
 - k8s (that's 'k' + 8 letters + 's')

WHAT IS KUBERNETES

- Kubernetes is a container cluster manager
- Manages containerized applications in a clustered environment
- It provides discovery across the cluster



- Kubelet (daemon on the worker node)
- Pods (collection of containers)
- Services (discovery mechanism between pods)
- Replication Controllers (replicated and monitored pods)

What is Kubernetes?

- Particularly suited for horizontally scaleable, stateless, or 'microservices' application architectures.
 - Does not mean others will not work or are ignored
- Additional functionality to make containers easier to use in a cluster (reachability and discovery).
- Kubernetes does NOT and will not expose all of the 'features' of the docker command line.

Kubernetes Key Words

- Master
- Node/Minion
- Pod
- Replication Controller
- Service
- Label
- Namespace

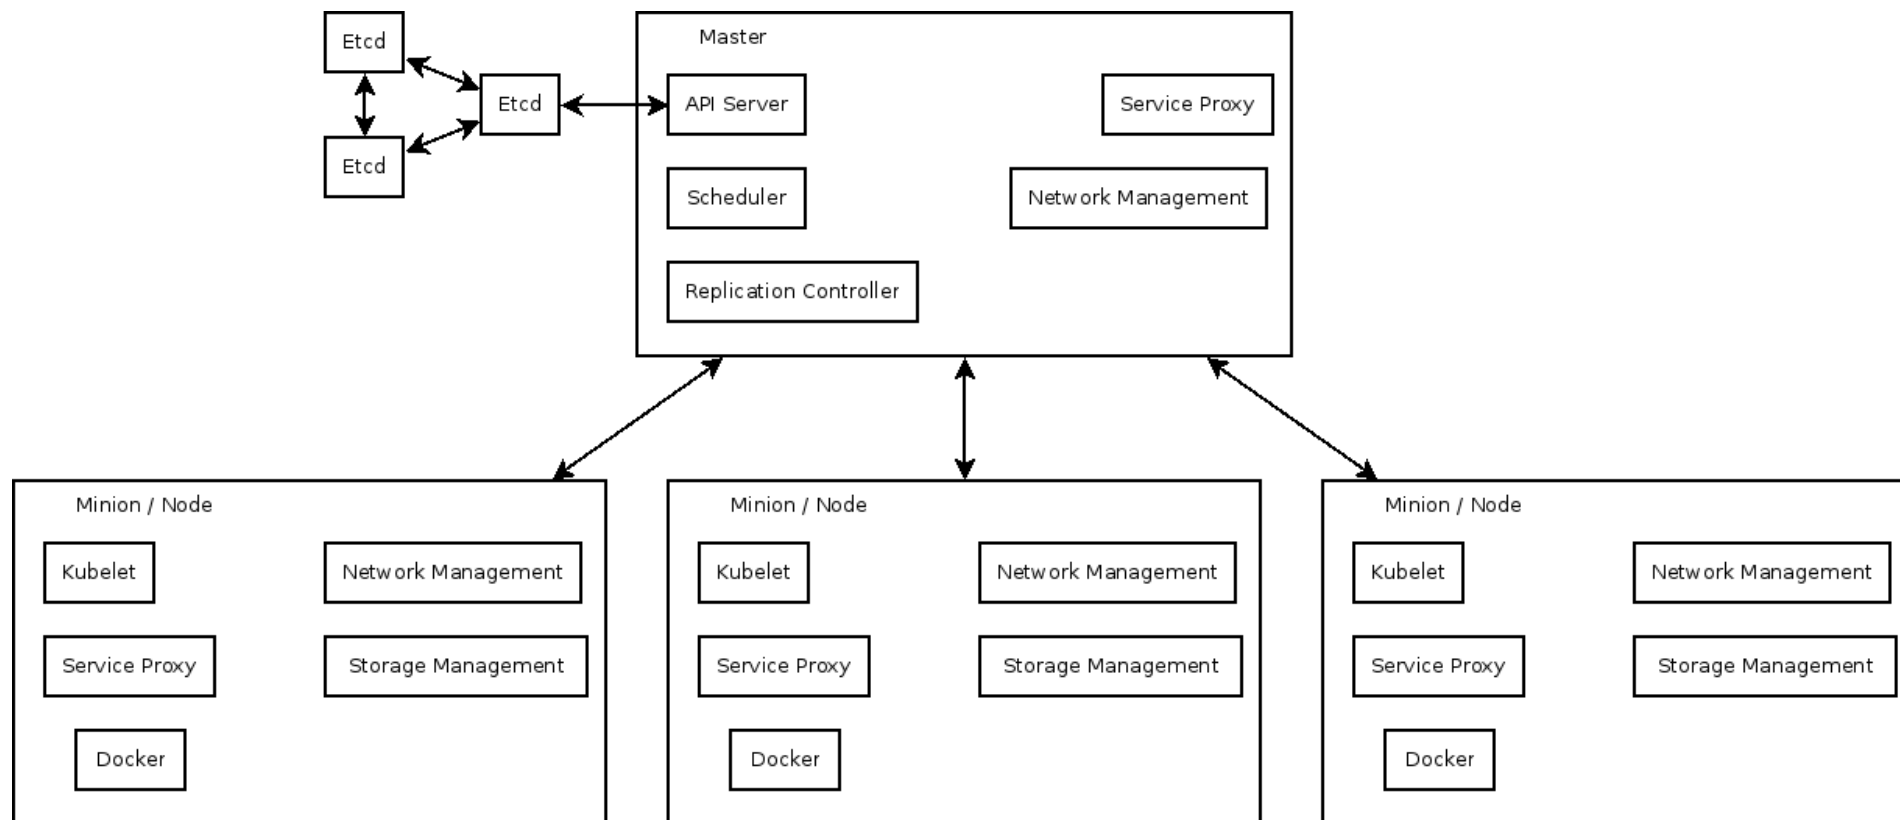
Master

- Typically consists of:
 - kube-apiserver
 - kube-scheduler
 - kube-controller-manager
 - etcd
- Might contain:
 - a network management utility

Node - Minion

- Typically consists of:
 - kubelet
 - kube-proxy
- Might contain:
 - a network management utility

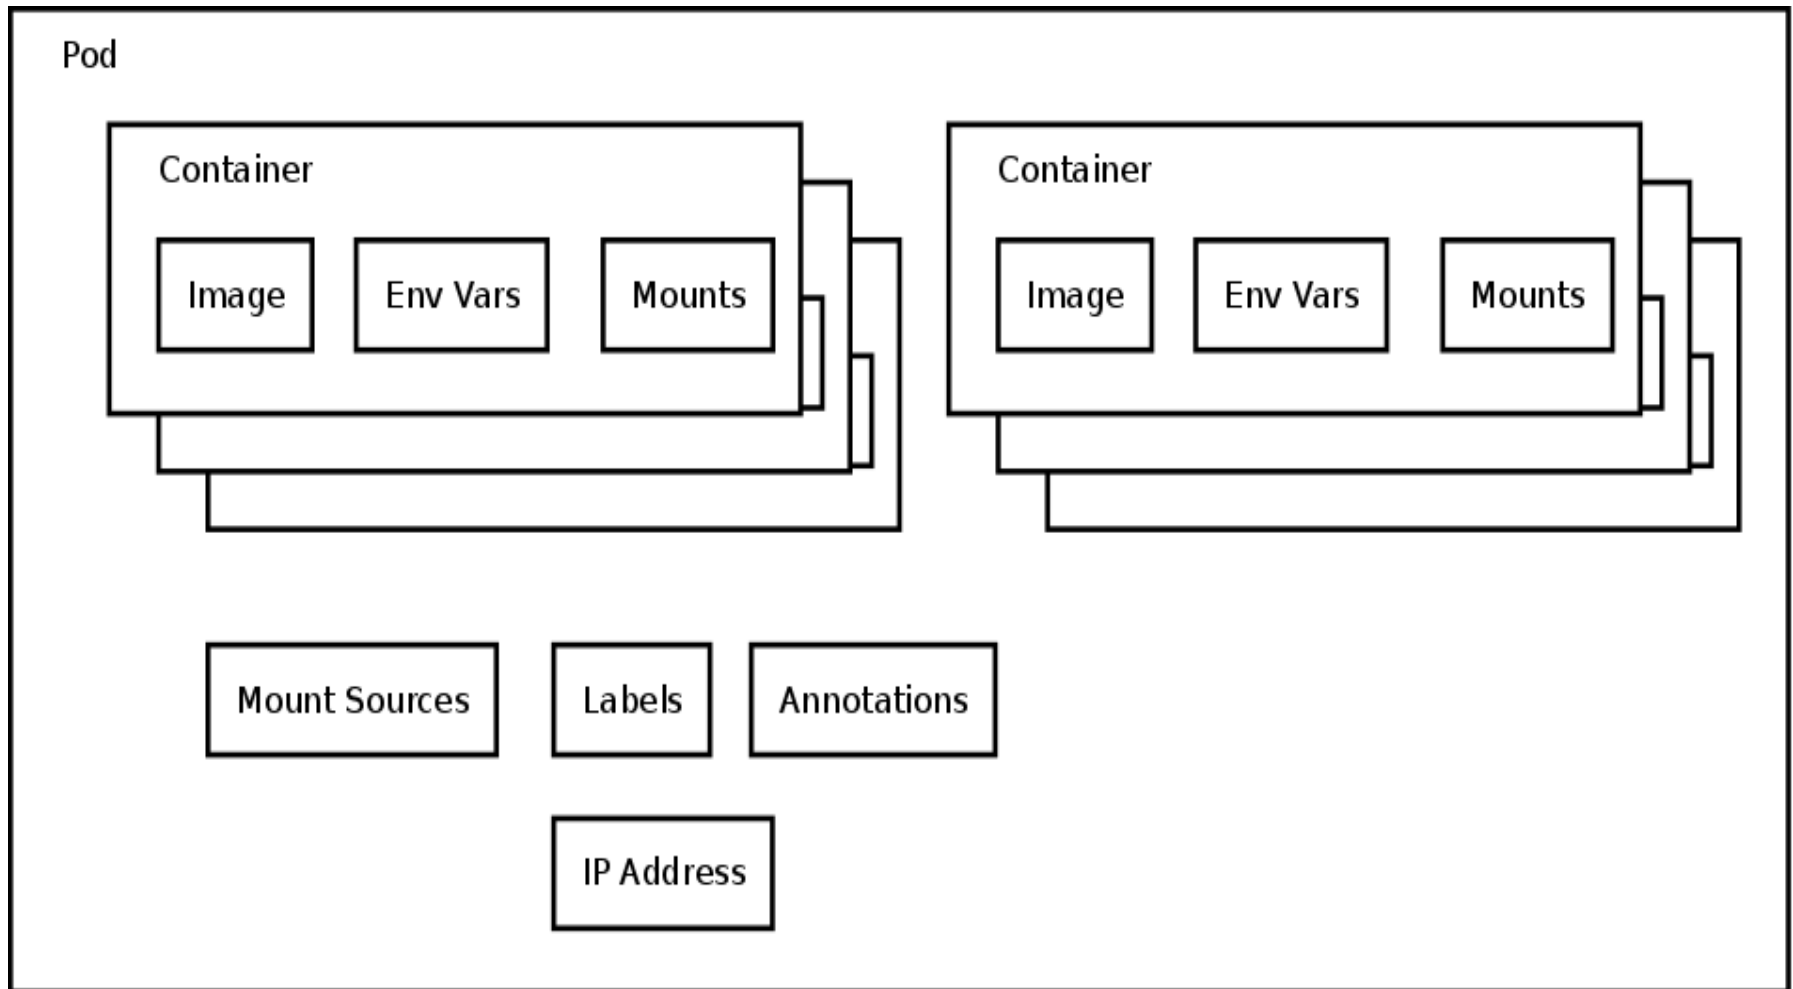
Architecture



Pod

- Single schedulable unit of work
 - Can not move between machines
- One or more containers
 - Shared network namespace
- Every pod gets an unique IP
 - Assigned by the container engine, not kube!
- Metadata about the container(s)
- Env vars – configuration for the container

Pod



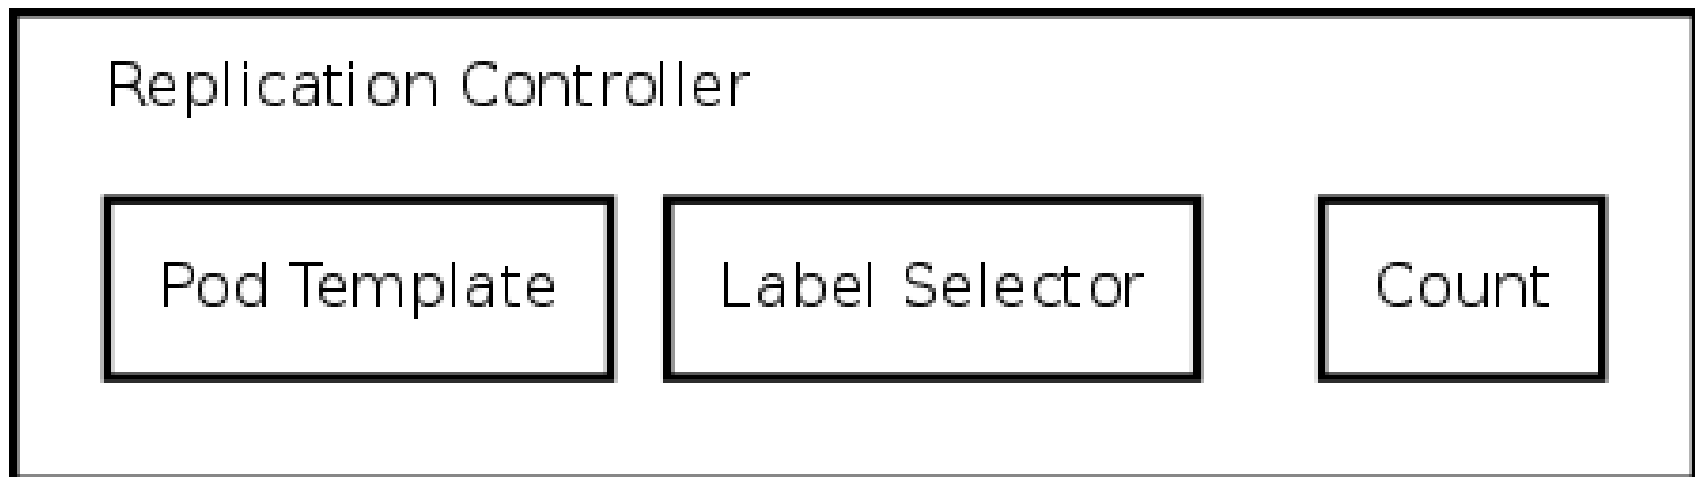
Pod - Example JSON

```
{ "apiVersion": "v1beta1", "id": "apache", "kind": "Pod", "namespace": "default",  
  "labels": { "name": "apache" },  
  "desiredState": { "manifest": {  
    "version": "v1beta1", "id": "apache", "volumes": null,  
    "containers": [{ "name": "my-fedora-apache", "image":  
"fedora/apache",  
    "ports": [{ "containerPort": 80, "hostPort": 80, "protocol": "TCP"  
}],  
    },  
    "restartPolicy": {" always": {} }  
  }, },  
}
```

Replication Controller

- Consists of
 - Pod template
 - Count
 - Label Selector
- Kube will try to keep \$count copies of pods matching the label selector running
- If too few copies are running the replication controller will start a new pod somewhere in the cluster

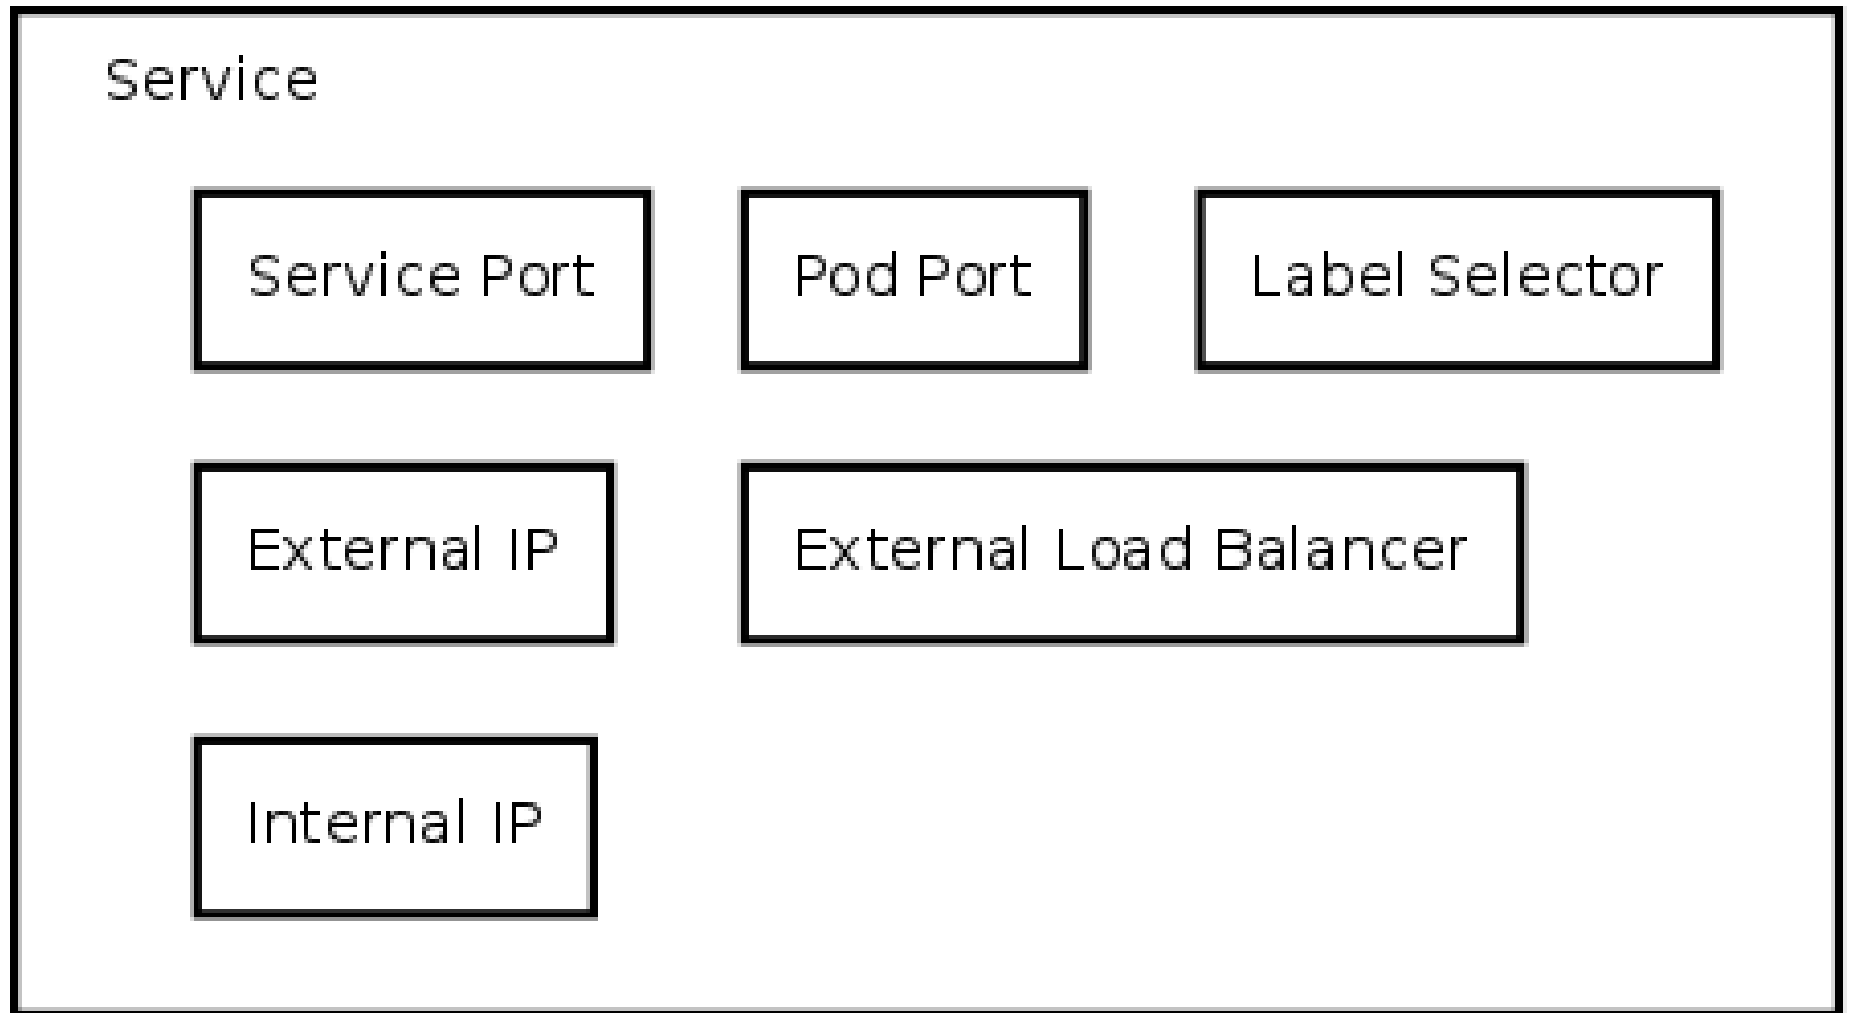
Replication Controller



Services

- Most every pod/replication controller will need a service. What's the point of a pod that doesn't provide some sort of service/useful work?
- How 'stuff' finds pods which could be anywhere?
 - Containers are started and stopped dynamically by kube, thus always changing IP addresses
- Define:
 - What port in the container
 - Labels on pods which should respond to this type of request

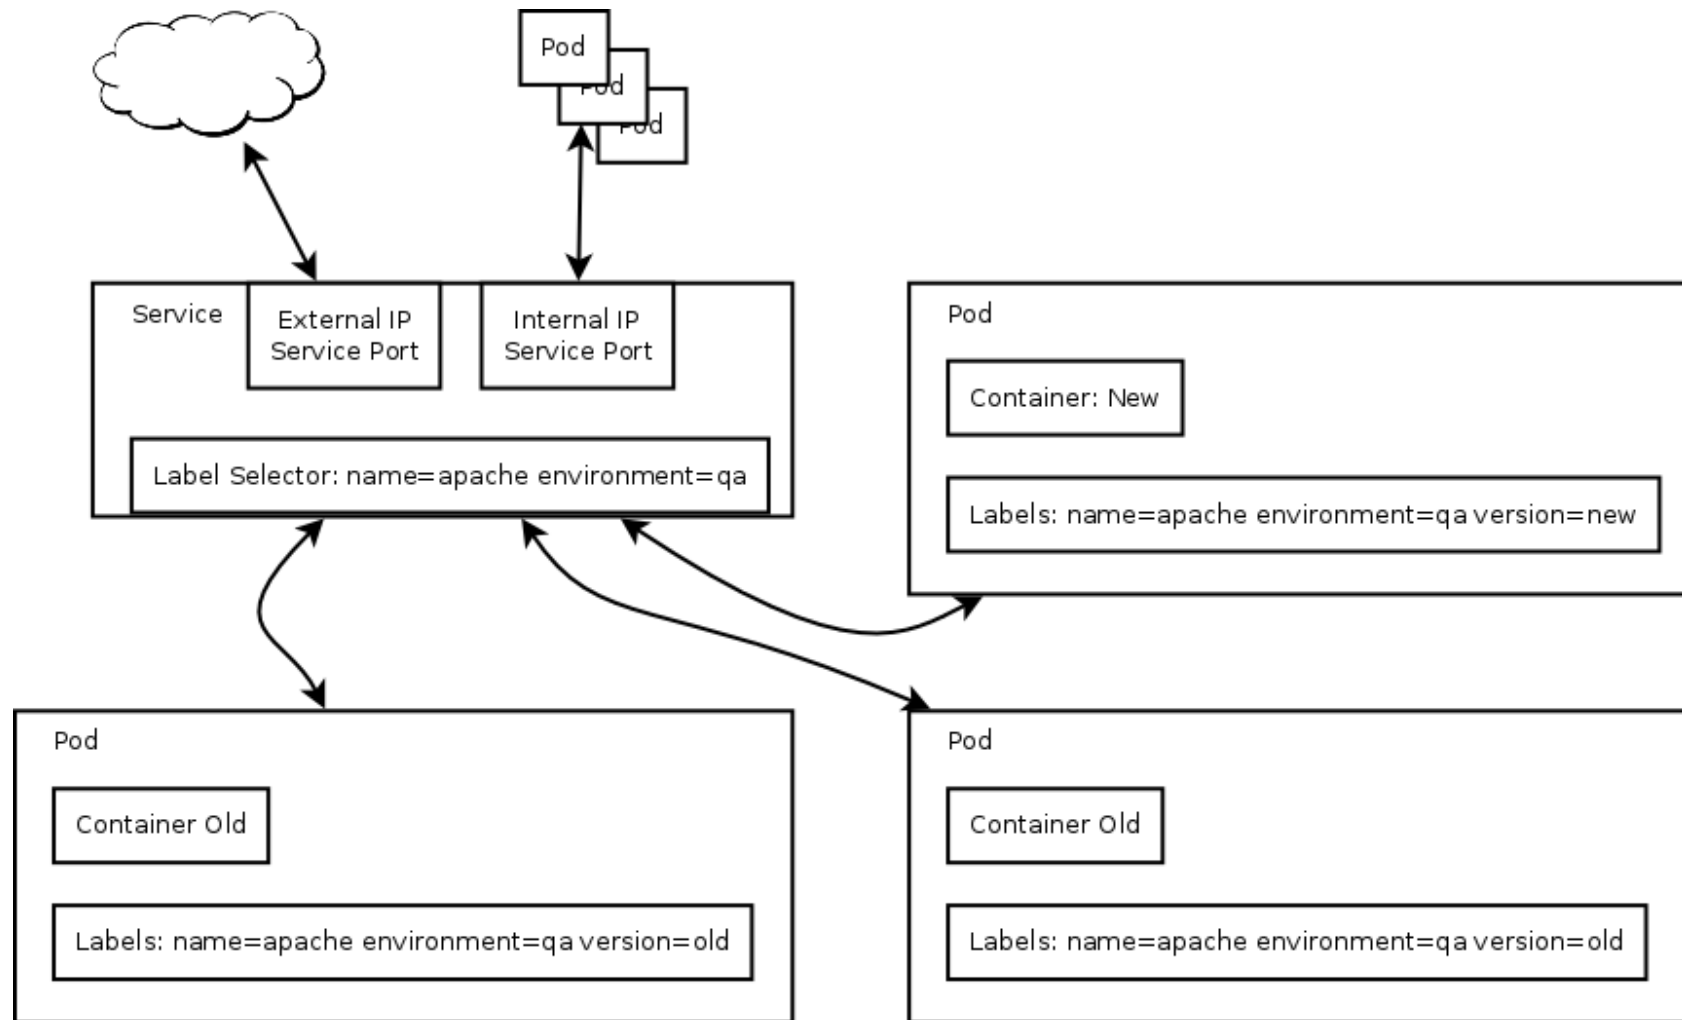
Services



Labels

- List of key=value pairs
- Attached to all objects
- Currently used in 2 main places
 - Matching pods to replication controllers
 - Matching pods to services
- Objects can be queried from the API server by label

Services and Labels



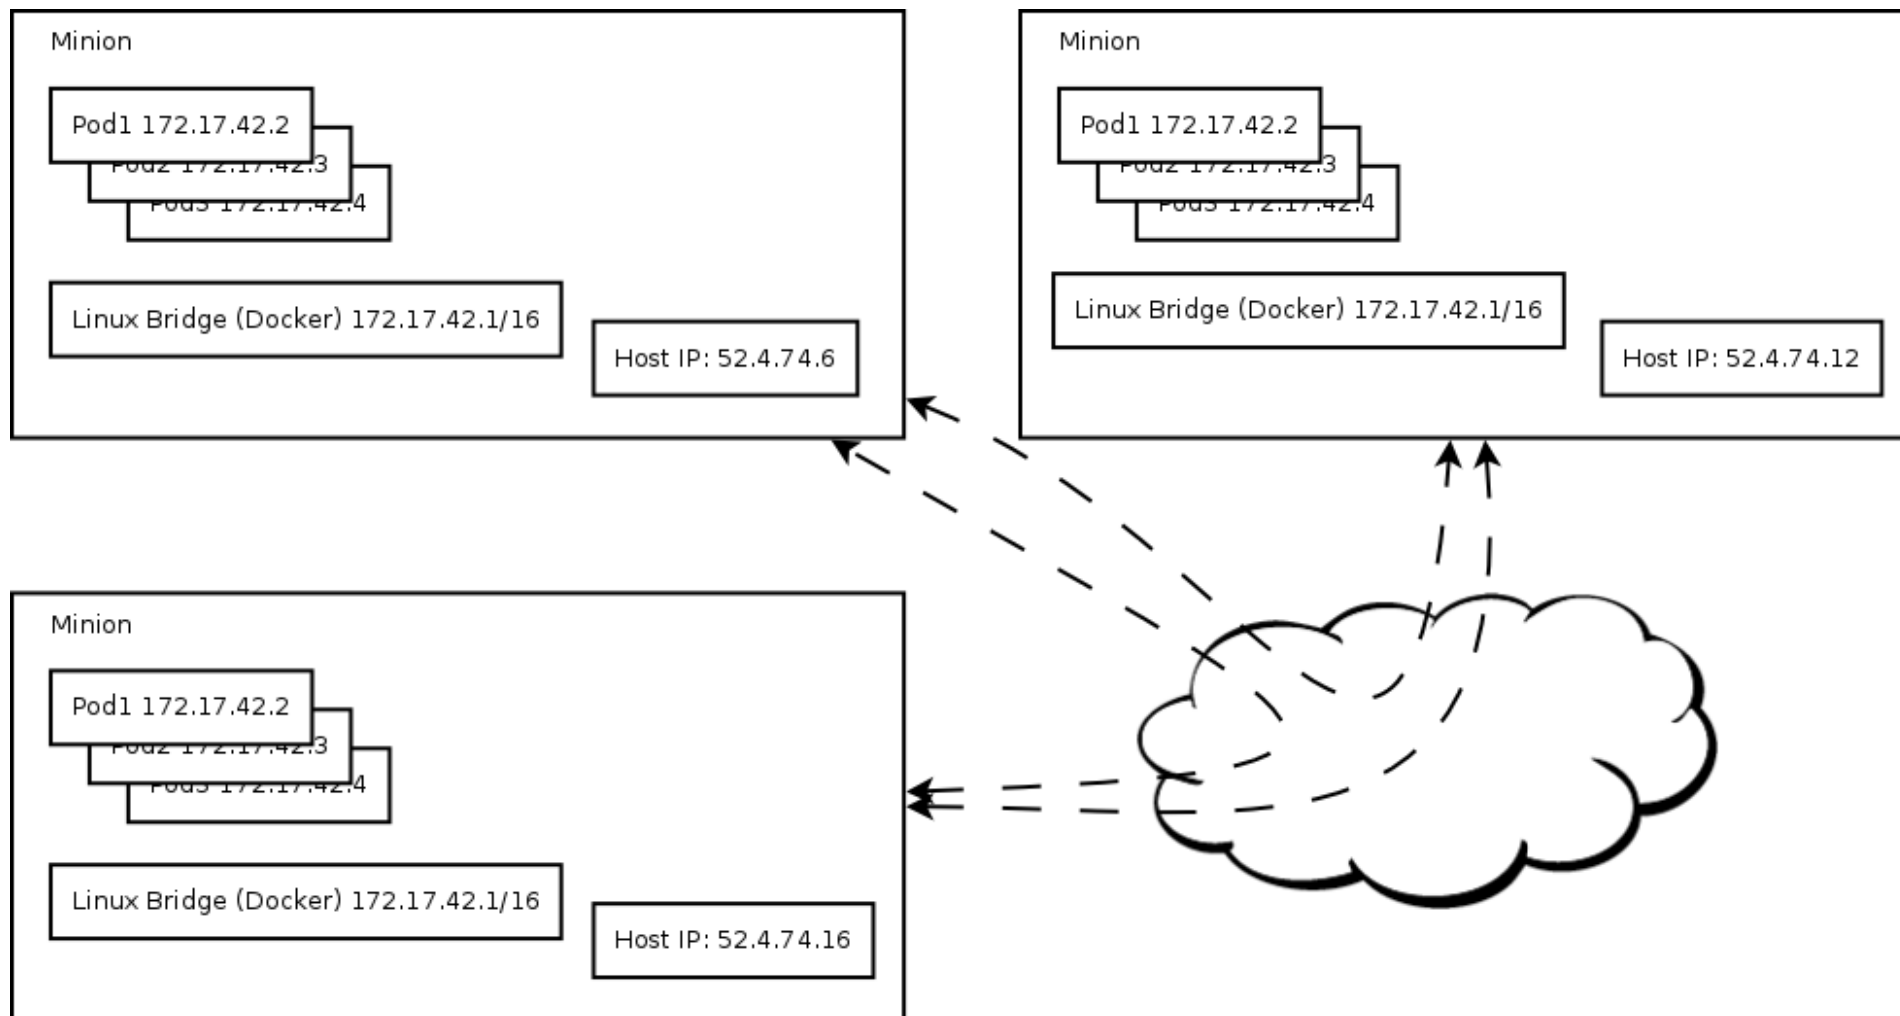
Namespace

- Attached to every object
- Pods in ns1 will not get service variable from ns2
- Users with permission to CRUD (create, read, update, delete) objects in ns1 may not have permissions to CRUD object in ns2
- The network is not segregated.
- Some people consider using a namespace per application. Some say a namespace per team or location.

Networking Setup

- Networking is a docker problem – not kube
 - Kube makes those problems apparent!
 - If any two docker containers on any two hosts can talk over IP, kube will just work.
- Docker looks so easy
 - 2 containers on one host can easily talk.
 - How to get to those containers from outside?
 - How to get to from one container on one host to a container on another?
- Networking is really hard!

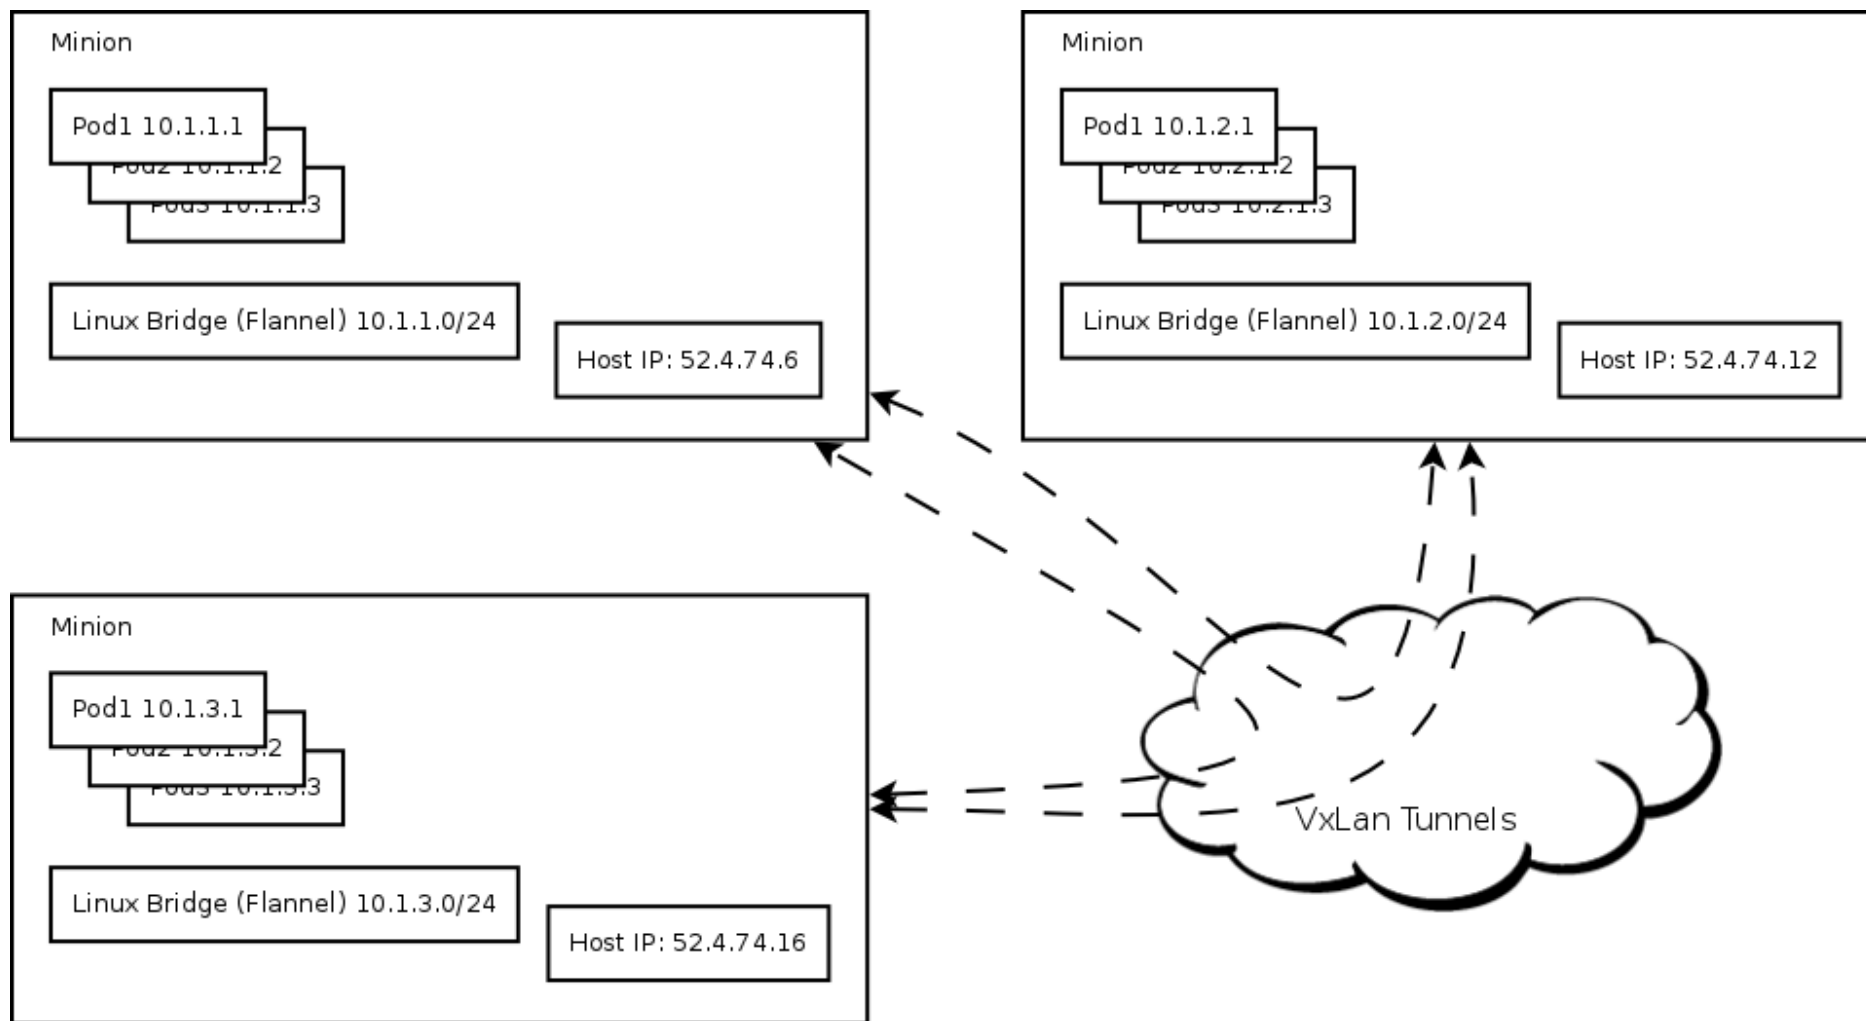
Networking Docker Out Of The Box



Networking Setup - Flannel

- Flannel
 - Super super easy configuration
 - Can create a vxlan overlay network
 - Can configure docker to launch pods in this overlay
 - Pods just work!
- There are many other solutions.
 - This one is easy.

Networking with an overlay network



Ressources

- Kubernetes

<http://kubernetes.io/>

- Project Atomic

<http://www.projectatomic.io/>

For questions please email

Thorsten Scherf <tscherf@redhat.com>

Credits go to Eric Paris for his excellent slidedeck.